

Retrieval-Augmented Generation (RAG) is a technique that combines retrieval from a knowledge base with text generation.

It works by first retrieving relevant documents from a vector database based on the user's query.

Then it feeds those documents as context to a large language model to generate an answer.

This approach reduces hallucination and improves answer accuracy.

Key components of RAG include: document chunking, embeddings, vector search, and LLM generation.